



Caption Consultation Brief

How do we make Saudi captions a creative director approves?

A complete, self-contained brief for an external AI/creative consultant.
Everything you need to advise us — assets, methods, what we tried, and the verdict.

The objective

We generate Arabic Instagram captions for Saudi brands. They are technically valid but **a Saudi creative director (Mohamed) rates ~83% of them weak or fail**. We want to know: **how do we consistently generate captions a Saudi creative director would approve?** This brief gives you our full context so you can advise — on methodology, prompting, architecture, or model strategy.

What we have (the assets)

Real data	6,888 verified Saudi IG posts; 3,077 captions with real likes
Templates	5,014 caption templates, tiered by real engagement
Brand brain	Voice, product names, signature hashtags for 23+ brands
Dialect rules	Saudi (Najdi/Hejazi) markers to use; Gulf/MSA words to avoid
Occasions	10 Saudi occasions w/ required words (Ramadan, founding day, etc.)
5 CD methodologies	Creative-director techniques (detailed below)
Quality gate	11 automated checks (dialect, product name, red lines, length)

The 5 Creative Director methodologies

Each caption is routed (by sector + occasion) to one of five creative techniques:

CD brain	Core technique
Firaasa Architect	Why-before-what — name the Saudi cultural truth first
Metaphor Architect	Build a full metaphor system, then a 'but wait' reveal
Authenticity Detective	Two-scene contrast: how people perform vs how they feel
Heritage Decoder	A double-meaning Arabic word as the structural key
Paradox Hunter	A counterintuitive flip; the product as the mechanism

How a caption is generated now

- 1. Route the brief (brand, product, occasion) to a CD methodology.
- 2. Pull 3-5 real high-engagement Saudi captions (templates) for that sector + occasion.
- 3. Build a prompt: brand voice + real templates + the CD technique + Saudi dialect rules + cultural red lines + length cap.
- 4. Generate with an LLM (gpt-4o-mini; we tested gpt-4o — no consistent gain).
- 5. Screen through the quality gate (dialect, product name, red lines, length).

The human verdict — what's actually wrong

Mohamed rated 30 captions blind. Result: **1 excellent, 4 good, 13 weak, 12 fail**. His notes revealed three concrete problems:

1. Too long — the #1 complaint (10 of 18 notes: 'too long', 'long and incoherent', 'boring'). Captions ran 300-400 characters. (*Fixed: hard length cap, now 116-172.*)

2. Cultural red lines violated — these are absolute:

- ممنوع ذكر إزالة أي ملابس أو الخمار/الحجاب
- ممنوع مشاهد السرير أو غرفة النوم
- ممنوع استغلال ضعف الناس أو هشاشتهم كخطاف

(*Fixed: hard-blocked in the quality gate.*)

3. The Authenticity Detective technique was the offender — its 'performance vs reality' two-scene structure generated long, intimate, inappropriate scenes (removing a veil, family on a bed). (*Fixed: constrained to public scenes.*)

Example — a FAIL (before fixes)

قهوة مختصة، لما تشعر بالجوع بعد يوم طويل من الصيام، تخيل اللحظة التي تفك فيها الخمار وتقول: أخيرا وقت الإفطار!
Mohamed's note: "we cannot use removing clothing in any caption." The AI judge scored this 6.8/10.

Example — an APPROVED (excellent)

في مشروع سكني، يستحضر التاريخ من عمق الأرض، حيث تبنى البيوت لتكون ملاذا وذكريات تروى. في يوم التأسيس، نحتفل بجذورنا.
Heritage Decoder — coherent heritage narrative, culturally safe, on-occasion.

A critical finding: AI cannot judge this

We built an AI 'methodology judge' (gpt-4o) to score captions. We validated it against Mohamed's ratings. **It scored essentially random** (rank correlation +0.08; 47% agreement). It gave ~6/10 to everything regardless of quality — including the culturally-forbidden captions. **Conclusion: automated AI scoring is not trustworthy for Saudi creative/cultural quality. Human review is the only ground truth.** Any solution you propose must be measurable against a human, not an AI judge.

What real Saudi engagement tells us

We analyzed 300 of the highest-engagement real Saudi captions through the 5 techniques:

Technique used	Share	Avg likes	vs avg
Counterintuitive flip	19%	6,593	+35%
Plain marketing (none)	9%	5,084	+4%
Double-meaning word	10%	4,920	+1%
Metaphor system	13%	4,912	+1%
Human-truth opening	48%	4,124	-15%

Insight: 91% of winning Saudi captions use a creative technique. The counterintuitive flip drives the most engagement but is underused; the emotional 'human-truth' opening is overused and underperforms.

Questions for you (the consultant)

- **Methodology:** Are these 5 CD techniques the right framework for short social captions, or are some (e.g. the two-scene Authenticity Detective) better suited to video than to a 150-char caption?
- **Prompting:** Given real templates + a named technique + dialect rules, what prompt architecture most reliably produces short, coherent, culturally-safe, genuinely-creative Saudi copy?
- **Length vs depth:** How do we keep a real creative technique (e.g. a double-meaning word) while staying under ~150 characters?
- **Model strategy:** We found gpt-4o didn't beat gpt-4o-mini with our prompt. Is that a prompt ceiling, or should we fine-tune a Saudi-dialect model?
- **The loop:** Human review is our only trusted signal but doesn't scale. How do we build a learning loop that improves quality without a reliable automated judge?

What success looks like

A batch of 30 fresh captions, rated blind by a Saudi creative director, where **the large majority are 'good' or 'excellent'** — short, culturally safe, on-brand, and genuinely creative. Today that number is 5 of 30. That is the gap to close.